

HOME CREDIT - CREDIT RISK MODEL STABILITY

채무 불이행 고객 예측

[Create](#)[Home](#)[Competitions](#)[Datasets](#)[Models](#)[Code](#)[Discussions](#)[Learn](#)[More](#)[Your Work](#)[VIEWED](#)[Home Credit - Credit ...](#)[Home Credit 2024 Sta...](#)[AI Mathematical Olym...](#)[Model Validation](#)[Your First Machine Le...](#)[EDITED](#)[Home Credit 2024 Sta...](#)[Exercise: Syntax, Varia...](#) Search

HOME CREDIT GROUP · FEATURED CODE COMPETITION · 23 DAYS TO GO

[Submit Prediction](#)

...

Home Credit - Credit Risk Model Stability

Create a model measured against feature stability over time

[Overview](#) [Data](#) [Code](#) [Models](#) [Discussion](#) [Leaderboard](#) [Rules](#) [Team](#) [Submissions](#)

Overview

The goal of this competition is to predict which clients are more likely to default on their loans. The evaluation will favor solutions that are stable over time.

Your participation may offer consumer finance providers a more reliable and longer-lasting way to assess a potential client's default risk.

Start

3 months ago

**Close**

23 days to go



Merger & Entry



Description

The absence of a credit history might mean a lot of things, including young age or a preference for cash. Without traditional data, someone with little to no credit history is likely to be denied. Consumer finance providers must accurately determine which clients can repay a loan and which cannot and data is key. If data science could help better predict one's repayment capabilities, loans might become more accessible to those who may benefit from them the most.

Currently, consumer finance providers use various statistical and machine learning methods to predict loan risk. These

Competition Host

Home Credit Group

**Prizes & Awards**

\$105,000

Awards Points & Medals

Participation

25,178 Entrants

3,680 Participants

2,916 Teams

44,444 Submissions

Tags[Tabular](#)[Banking](#)[Custom Metric](#)**Table of Contents**[Overview](#)[Description](#)[Evaluation](#)

HOME CREDIT GROUP



**HOME
CREDIT**

The logo consists of the words "HOME" and "CREDIT" stacked vertically. "HOME" is on top, with its letters slightly slanted. "CREDIT" is below it, also with its letters slightly slanted. Both words are in a bold, red, sans-serif font.


다국적 금융 회사

신용도가 낮은 소비자에게 대출 제공

라이브러리 및 데이터 불러오기

```
import polars as pl
import numpy as np
import pandas as pd
import lightgbm as lgb
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score

dataPath = "C://Users//noon9//Untitled Folder//home-credit-credit-risk-model-stability//"
```



대용량 데이터를 처리하는 라이브러리

TRAIN_STATIC

| case_id | actualdp tolerance_344P | amtinstpaidbefduel24m_4187115A | annuity_780A | annuitynextmonth_57A | applicationcnt_361L | applications30d_658L | applicati |
|---------|-------------------------|--------------------------------|--------------|----------------------|---------------------|----------------------|-----------|
| i64 | f64 | str | f64 | f64 | f64 | f64 | f64 |
| 0 | null | null | 1917.6 | 0.0 | 0.0 | 0.0 | 0.0 |
| 1 | null | null | 3134.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 2 | null | null | 4937.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 3 | null | null | 4643.6 | 0.0 | 0.0 | 0.0 | 1.0 |
| 4 | null | null | 3390.2 | 0.0 | 0.0 | 0.0 | 1.0 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 2651088 | 0.0 | "117624.79" | 3045.0 | 4488.0 | 0.0 | 0.0 | 0.0 |
| 2651089 | 0.0 | "134887.4" | 1200.0 | 4382.8003 | 0.0 | 0.0 | 0.0 |
| 2651090 | 0.0 | "69186.62" | 6000.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 2651091 | 0.0 | "117331.0" | 11565.4 | 13781.2 | 0.0 | 0.0 | 0.0 |
| 2651092 | 0.0 | "44442.2" | 3045.0 | 3000.0 | 0.0 | 0.0 | 1.0 |

TEST_STATIC

| | | | | | | | |
|---------|-------------------------|--------------------------------|--------------|----------------------|---------------------|----------------------|--------------------------|
| case_id | actualdpdtolerance_344P | amtinstpaidbefdue124m_4187115A | annuity_780A | annuitynextmonth_57A | applicationcnt_361L | applications30d_658L | applicationnextmonth_57A |
| i64 | f64 | f64 | f64 | f64 | f64 | f64 | f64 |
| 57543 | 0.0 | 191767.36 | 3674.6 | 1218.2001 | 0.0 | 0.0 | 0.0 |
| 57551 | 0.0 | 71036.4 | 2844.6 | 0.0 | 0.0 | 1.0 | 0.0 |
| 57552 | 0.0 | 183992.0 | 6298.8003 | 12155.4 | 0.0 | 0.0 | 0.0 |
| 57569 | 0.0 | 0.0 | 4682.6 | 0.0 | 0.0 | 1.0 | 0.0 |
| 57630 | 0.0 | 0.0 | 8905.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 57631 | 0.0 | null | 2540.6 | 0.0 | 0.0 | 0.0 | 0.0 |
| 57632 | 0.0 | 63647.402 | 4732.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 57635 | 0.0 | null | 1167.4 | 0.0 | 0.0 | 0.0 | 0.0 |
| 57637 | 0.0 | 43677.184 | 4300.4 | 0.0 | 0.0 | 0.0 | 0.0 |
| 57639 | 0.0 | 142333.14 | 12599.601 | 0.0 | 0.0 | 0.0 | 0.0 |

case_id - 신용 건별 고유 식별자

date_decision - 대출승인이 결정됐던 날

Target - 고객이 특정 신용건(대출)에 대한 채무불이행 여부

WEEK_NUM - 주 단위로 데이터를 집계하는 데 사용되는 주 번호

MONTH - 월 단위 집계

num_group - 과거 기록을 위한 인덱싱

 **P** - Transform DPD (Days past due)

 **M** - Masking categories

 **A** - Transform amount

 **D** - Transform date

 **T** - Unspecified Transform

 **L** - Unspecified Transform

. pmts_month_158T : 활성계약에 대한 지불 내역

pmts_month_706T : 종료계약에 대한 지불 내역

dateofcredstart_18ID : 신용계약의 시작날짜

데이터 타입 설정 함수 (POLARS)

```
def set_table_dtypes(df: pl.DataFrame) -> pl.DataFrame:
    for col in df.columns:
        if col[-1] in ("P", "A"):
            df = df.with_columns(pl.col(col).cast(pl.Float64).alias(col))
    return df
```

컬럼 이름의 마지막 문자가 'P' 또는 'A'인 경우, **Float64**로 변환

변환 함수 (PANDAS)

```
def convert_strings(df: pd.DataFrame) -> pd.DataFrame:
    for col in df.columns:
        if df[col].dtype.name in ['object', 'string']:

            df[col] = df[col].astype("string").astype('category')
            current_categories = df[col].cat.categories
            new_categories = current_categories.to_list() + ["Unknown"]
            new_dtype = pd.CategoricalDtype(categories=new_categories, ordered=True)
            df[col] = df[col].astype(new_dtype)
    return df
```

문자열 컬럼을 범주형(카테고리)데이터로 변환

"Unknown" 카테고리 추가

트레인 데이터 로드 및 전처리

```
train_basetable = pl.read_csv(dataPath + "csv_files//train//train_base.csv")
train_static = pl.concat(
    [
        pl.read_csv(dataPath + "csv_files//train//train_static_0_0.csv").pipe(set_table_dtypes),
        pl.read_csv(dataPath + "csv_files//train//train_static_0_1.csv").pipe(set_table_dtypes),
    ],
    how="vertical_relaxed",
)
train_static_cb = pl.read_csv(dataPath + "csv_files//train//train_static_cb_0.csv").pipe(set_table_dtypes)
train_person_1 = pl.read_csv(dataPath + "csv_files//train//train_person_1.csv").pipe(set_table_dtypes)
train_credit_bureau_b_2 = pl.read_csv(dataPath + "csv_files//train//train_credit_bureau_b_2.csv").pipe(set_table_dtypes)
```

TRAIN_PERSON_I

| train_person_1 | | | | | | | | |
|------------------------|--------------|---------------|---------------|-----------------------|--------------------------|--------------------------|-----------------------|-----|
| shape: (2_973_991, 37) | | | | | | | | |
| case_id | birth_259D | birthdate_87D | childnum_185L | contaddr_district_15M | contaddr_matchlist_1032L | contaddr_smempladdr_334L | contaddr_zipcode_807M | ed |
| 0 | "1986-07-01" | null | null | "P88_18_84" | false | false | "P167_100_165" | |
| 0 | null | null | null | "a55475b1" | null | null | "a55475b1" | |
| 0 | null | null | null | "a55475b1" | null | null | "a55475b1" | |
| 0 | null | null | null | "a55475b1" | null | null | "a55475b1" | |
| 1 | "1957-08-01" | null | null | "P103_93_94" | false | false | "P176_37_166" | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 2703451 | null | null | null | "a55475b1" | null | null | "a55475b1" | |
| 2703452 | "1977-08-01" | null | null | "P133_44_167" | false | false | "P59_150_74" | |
| 2703453 | "1950-02-01" | null | null | "P123_6_84" | false | false | "P46_103_143" | |
| 2703453 | null | null | null | "a55475b1" | null | null | "a55475b1" | |
| 2703454 | "1948-04-01" | null | null | "P48_127_19" | false | false | "P78_144_175" | |

TRAIN_PERSON_1 테이블 그룹화 및 집계

```
train_person_1_feats_1 = train_person_1.groupby("case_id").agg(  
    pl.col("mainoccupationinc_384A").max().alias("mainoccupationinc_384A_max"),  
    (pl.col("incometype_1044T") == "SELFEMPLOYED").max().alias("mainoccupationinc_384A_any_selfemployed")  
)  
train_person_1_feats_1
```

shape: (1_526_659, 3)

| case_id | mainoccupationinc_384A_max | mainoccupationinc_384A_any_selfemployed |
|---------|----------------------------|---|
| i64 | f64 | bool |
| 2535791 | 72000.0 | false |
| 1005921 | 22000.0 | false |
| 1301441 | 40000.0 | false |
| 1357092 | 80000.0 | false |
| 785388 | 77000.0 | false |
| ... | ... | ... |
| 1768097 | 34000.0 | false |
| 909602 | 30000.0 | false |
| 1412642 | 60000.0 | false |
| 775017 | 40000.0 | false |
| 1798099 | 40000.0 | false |

TRAIN_PERSON_I 테이블 필터링 및 열 선택

```
train_person_1_feats_2 = train_person_1.select(["case_id", "num_group1", "housetype_905L"]).filter(  
    pl.col("num_group1") == 0  
).drop("num_group1").rename({"housetype_905L": "person_housetype"})  
train_person_1_feats_2
```

```
shape: (1_526_659, 2)
```

| case_id | person_housetype |
|---------|------------------|
| 0 | null |
| 1 | null |
| 2 | null |
| 3 | null |
| 4 | null |
| ... | ... |
| 2703450 | "OWNED" |
| 2703451 | null |
| 2703452 | null |
| 2703453 | null |
| 2703454 | null |

TRAIN_CREDIT_BUREAU_B_2

```
train_credit_bureau_b_2
```

```
shape: (1_286_755, 6)
```

| case_id | num_group1 | num_group2 | pmts_date_1107D | pmts_dpvalue_108P | pmts_pmtsoverdue_635A |
|---------|------------|------------|-----------------|-------------------|-----------------------|
| i64 | i64 | i64 | str | f64 | f64 |
| 467 | 0 | 0 | "2018-11-15" | null | null |
| 467 | 0 | 1 | "2018-12-15" | null | null |
| 467 | 1 | 0 | "2018-12-15" | null | null |
| 467 | 2 | 0 | "2016-10-15" | 0.0 | 0.0 |
| 467 | 2 | 1 | "2016-11-15" | 0.0 | 0.0 |
| ... | ... | ... | ... | ... | ... |
| 2703436 | 1 | 31 | "2020-05-15" | 0.0 | 0.0 |
| 2703436 | 1 | 32 | "2020-06-15" | 0.0 | 0.0 |
| 2703436 | 1 | 33 | "2020-07-15" | 0.0 | 0.0 |
| 2703436 | 1 | 34 | "2020-08-15" | 0.0 | 0.0 |
| 2703436 | 1 | 35 | "2020-09-15" | 0.0 | 0.0 |

TRAIN_CREDIT_BUREAU_B_2 테이블 그룹화 및 집계

```
train_credit_bureau_b_2_feats = train_credit_bureau_b_2.groupby("case_id").agg(  
    pl.col("pmts_pmtsoverdue_635A").max().alias("pmts_pmtsoverdue_635A_max"),  
    (pl.col("pmts_dpvalue_108P") > 31).max().alias("pmts_dpvalue_108P_over31"))  
)  
train_credit_bureau_b_2_feats  
shape: (36_447, 3)
```

| case_id | pmts_pmtsoverdue_635A_max | pmts_dpvalue_108P_over31 |
|---------|---------------------------|--------------------------|
| i64 | f64 | bool |
| 53621 | 0.0 | false |
| 223684 | 0.0 | false |
| 1713554 | 0.0 | false |
| 1708984 | 0.0 | false |
| 744685 | 0.0 | false |
| ... | ... | ... |
| 882429 | 0.0 | false |
| 1904001 | 0.8 | true |
| 1938255 | 0.4 | true |
| 1447979 | 0.0 | false |
| 914271 | 0.0 | false |

선택된 A 및 M 타입 열 추출

```
selected_static_cols = []
for col in train_static.columns:
    if col[-1] in ("A", "M"):
        selected_static_cols.append(col)
print(selected_static_cols)

selected_static_cb_cols = []
for col in train_static_cb.columns:
    if col[-1] in ("A", "M"):
        selected_static_cb_cols.append(col)
print(selected_static_cb_cols)

['amtinstpaidbefdue124m_4187115A', 'annuity_780A', 'annuitynextmonth_57A', 'avginstalllast24m_3658937A', 'avgnamtstart24m_4525187A', 'avgoutstandbalance16m_4187114A', 'avgpmtlast12m_4525200A', 'credamount_770A', 'currdebt_22A', 'currdebtcreditperange_828A', 'disbursedcredamount_1113A', 'downpmt_116A', 'inittransactionamount_650A', 'lastapprcommoditycat_1041M', 'lastapprcommoditytypec_5251766M', 'lastapprcredamount_781A', 'lastcancelreason_561M', 'lastotherinc_902A', 'lastotherinsexpense_631A', 'lastrejectcommoditycat_161M', 'lastrejectcommoditytypec_5251769M', 'lastrejectcredamount_222A', 'lastrejectreason_759M', 'lastrejectreasonclient_4145040M', 'maininc_215A', 'maxannuity_159A', 'maxannuity_4075009A', 'maxdebt4_972A', 'maxinstalllast24m_3658928A', 'maxlnamtstart6m_4525199A', 'maxoutstandbalance12m_4187113A', 'maxpmtlast3m_4525190A', 'previouscontdistrict_112M', 'price_1097A', 'sumoutstandtotal_3546847A', 'sumoutstandtotalest_4493215A', 'totaldebt_9A', 'totalsettled_863A', 'tostinstalllast1m_4525188A']
['description_5085714M', 'education_1103M', 'education_88M', 'maritalst_385M', 'maritalst_893M', 'pmtaverage_3A', 'pmtaverage_4527227A', 'pmtaverage_4955615A', 'pmtssum_45A']
```

모든 테이블 결합

```
data = train_basetable.join(
    train_static.select(["case_id"]+selected_static_cols), how="left", on="case_id"
).join(
    train_static_cb.select(["case_id"]+selected_static_cb_cols), how="left", on="case_id"
).join(
    train_person_1_feats_1, how="left", on="case_id"
).join(
    train_person_1_feats_2, how="left", on="case_id"
).join(
    train_credit_bureau_b_2_feats, how="left", on="case_id"
)
data.head()
```

shape: (5, 58)

| case_id | date_decision | MONTH | WEEK_NUM | target | amtinstpaidbefduel24m_4187115A | annuity_780A | annuitynextmonth_57A | avginstalllast24m_3658937A | avgli |
|---------|---------------|--------|----------|--------|--------------------------------|--------------|----------------------|----------------------------|-------|
| i64 | str | i64 | i64 | i64 | f64 | f64 | f64 | f64 | f64 |
| 0 | "2019-01-03" | 201901 | 0 | 0 | null | 1917.6 | 0.0 | null | |
| 1 | "2019-01-03" | 201901 | 0 | 0 | null | 3134.0 | 0.0 | null | |
| 2 | "2019-01-04" | 201901 | 0 | 0 | null | 4937.0 | 0.0 | null | |
| 3 | "2019-01-03" | 201901 | 0 | 0 | null | 4643.6 | 0.0 | null | |
| 4 | "2019-01-04" | 201901 | 0 | 1 | null | 3390.2 | 0.0 | null | |

테스트 데이터 로드 및 전처리

트레인 데이터와 동일한 처리 방식

```
test_basetable = pl.read_csv(dataPath + "csv_files//test//test_base.csv")
test_static = pl.concat(
    [
        pl.read_csv(dataPath + "csv_files//test//test_static_0_0.csv").pipe(set_table_dtypes),
        pl.read_csv(dataPath + "csv_files//test//test_static_0_1.csv").pipe(set_table_dtypes),
        pl.read_csv(dataPath + "csv_files//test//test_static_0_2.csv").pipe(set_table_dtypes),
    ],
    how="vertical_relaxed",
)
test_static_cb = pl.read_csv(dataPath + "csv_files//test//test_static_cb_0.csv").pipe(set_table_dtypes)
test_person_1 = pl.read_csv(dataPath + "csv_files//test//test_person_1.csv").pipe(set_table_dtypes)
test_credit_bureau_b_2 = pl.read_csv(dataPath + "csv_files//test//test_credit_bureau_b_2.csv").pipe(set_table_dtypes)
```

테스트 데이터 로드 및 전처리

트레인 데이터와 동일한 처리 방식

```
# test_person_1 데이터를 그룹화 및 카운트
test_person_1_feats_1 = test_person_1.group_by("case_id").agg(
    pl.col("mainoccupationinc_384A").max().alias("mainoccupationinc_384A_max"),
    (pl.col("incometype_1044T") == "SELFEMPLOYED").max().alias("mainoccupationinc_384A_any_selfemployed")
)

# test_person_1 데이터를 필터링 및 열 선택
test_person_1_feats_2 = test_person_1.select(["case_id", "num_group1", "housetype_905L"]).filter(
    pl.col("num_group1") == 0
).drop("num_group1").rename({"housetype_905L": "person_housetype"})

# test_credit_bureau_b_2 데이터를 그룹화 및 카운트
test_credit_bureau_b_2_feats = test_credit_bureau_b_2.group_by("case_id").agg(
    pl.col("pmts_pmtsoverdue_635A").max().alias("pmts_pmtsoverdue_635A_max"),
    (pl.col("pmts_dpdpvalue_108P") > 31).max().alias("pmts_dpdpvalue_108P_over31")
)

# 모든 데이터 결합
data_submission = test_basetable.join(
    test_static.select(["case_id"] + selected_static_cols), how="left", on="case_id"
).join(
    test_static_cb.select(["case_id"] + selected_static_cb_cols), how="left", on="case_id"
).join(
    test_person_1_feats_1, how="left", on="case_id"
).join(
    test_person_1_feats_2, how="left", on="case_id"
).join(
    test_credit_bureau_b_2_feats, how="left", on="case_id"
)
```

data_submission

shape: (10, 57)

| case_id | date_decision | MONTH | WEEK_NUM | amtinstpaidbefduel24m_4187115A | annuity_780A | annuitynextmonth_57A | avginstalllast24m_3658937A | avglnamtsta |
|---------|---------------|--------|----------|--------------------------------|--------------|----------------------|----------------------------|-------------|
| i64 | str | i64 | i64 | f64 | f64 | f64 | f64 | f64 |
| 57543 | "2021-05-14" | 202201 | 100 | 191767.36 | 3674.6 | 1218.2001 | 16049.4 | |
| 57549 | "2022-01-17" | 202201 | 100 | 129704.4 | 5742.6 | 3546.6 | 32426.201 | |
| 57551 | "2020-11-27" | 202201 | 100 | 71036.4 | 2844.6 | 0.0 | 8357.2 | |
| 57552 | "2020-11-27" | 202201 | 100 | 183992.0 | 6298.8003 | 12155.4 | 7440.4 | |
| 57569 | "2021-12-20" | 202201 | 100 | 0.0 | 4682.6 | 0.0 | null | |
| 57630 | "2021-03-16" | 202201 | 100 | 0.0 | 8905.0 | 0.0 | null | |
| 57631 | "2022-06-04" | 202201 | 100 | null | 2540.6 | 0.0 | null | |
| 57632 | "2022-02-05" | 202201 | 100 | 63647.402 | 4732.0 | 0.0 | 3536.0 | |
| 57633 | "2022-01-25" | 202201 | 100 | null | 8273.0 | 0.0 | null | |
| 57634 | "2021-01-27" | 202201 | 100 | 39948.8 | 1165.8 | 0.0 | 3994.8 | |

데이터 세트를 훈련, 검증 및 테스트 세트로 분할

```
case_ids = data["case_id"].unique().shuffle(seed=1)
case_ids_train, case_ids_test = train_test_split(case_ids, train_size=0.6, random_state=1)
case_ids_valid, case_ids_test = train_test_split(case_ids_test, train_size=0.5, random_state=1)

cols_pred = []
for col in data.columns:
    if col[-1].isupper() and col[:-1].islower():
        cols_pred.append(col)
```

POLARS 데이터를 PANDAS 데이터로 변환

```
def from_polars_to_pandas(case_ids: pl.DataFrame) -> pl.DataFrame:
    return (
        data.filter(pl.col("case_id").is_in(case_ids))[[ "case_id", "WEEK_NUM", "target"]].to_pandas(),
        data.filter(pl.col("case_id").is_in(case_ids))[cols_pred].to_pandas(),
        data.filter(pl.col("case_id").is_in(case_ids))["target"].to_pandas()
    )
```

```
base_train, X_train, y_train = from_polars_to_pandas(case_ids_train)
base_valid, X_valid, y_valid = from_polars_to_pandas(case_ids_valid)
base_test, X_test, y_test = from_polars_to_pandas(case_ids_test)
```

```
for df in [X_train, X_valid, X_test]:
    df = convert_strings(df)
```

훈련, 검증 및 테스트 세트의 크기

```
print(f"Train: {X_train.shape}")
print(f"Valid: {X_valid.shape}")
print(f"Test: {X_test.shape}")

Train: (915995, 48)
Valid: (305332, 48)
Test: (305332, 48)
```

LIGHTGBM 모델 훈련

```
lgb_train = lgb.Dataset(X_train, label=y_train)
lgb_valid = lgb.Dataset(X_valid, label=y_valid, reference=lgb_train)

params = {
    "boosting_type": "gbdt",
    "objective": "binary",
    "metric": "auc",
    "max_depth": 3,
    "num_leaves": 31,
    "learning_rate": 0.05,
    "feature_fraction": 0.9,
    "bagging_fraction": 0.8,
    "bagging_freq": 5,
    "n_estimators": 1000,
    "verbose": -1,
}



gbm = lgb.train(
    params,
    lgb_train,
    valid_sets=lgb_valid,
    callbacks=[lgb.log_evaluation(50), lgb.early_stopping(10)]
)
```

LIGHTGBM



리프 기준 분할

LEAF-WISE



LIGHTGBM 파라미터

| 파라미터 | default | 설명 |
|------------------|------------|--|
| num_iterations | 100 | 반복 수행하려는 트리의 개수 (너무 크면 오버피팅 발생) |
| objective | regression | 수치예측이면 regression, 이진분류이면 binary |
| learning_rate | 0.1 | 부스팅 스텝 반복할 때 학습률, 0~1 사이의 값 |
| max_depth | 1 | 트리의 깊이 |
| min_data_in_leaf | 20 | 한 리프의 최소 데이터 수 (decision tree의 min_sample_leaf와 동일, 오버피팅 제어) |
| num_leaves | 31 | 하나의 트리가 가질 수 있는 최대 리프 개수 |
| boosting | gbdt | 부스팅 방법 (gbdt: Gradient Boosting DecisionTree / rf: RandomForest) |
| bagging_fraction | 1.0 | 데이터 샘플링 비율, 오버피팅 제어 |
| feature_fraction | 1.0 | 개별 트리 학습 시 무작위로 선택하는 feature의 비율 |
| lambda_l1 | 0.0 | L1 regulation 제어 |
| lambda_l2 | 0.0 | L2 regulation 제어 |
| metric | "" | 성능평가를 어떤 것으로 할 것인지 (auc, l1, l2 등) |

LIGHTGBM 모델 훈련

```
Training until validation scores don't improve for 10 rounds
[50]    valid_0's auc: 0.705963
[100]   valid_0's auc: 0.724362
[150]   valid_0's auc: 0.731423
[200]   valid_0's auc: 0.735874
[250]   valid_0's auc: 0.739009
[300]   valid_0's auc: 0.740965
[350]   valid_0's auc: 0.742924
[400]   valid_0's auc: 0.744582
[450]   valid_0's auc: 0.745977
[500]   valid_0's auc: 0.747033
[550]   valid_0's auc: 0.747877
[600]   valid_0's auc: 0.749039
[650]   valid_0's auc: 0.750087
[700]   valid_0's auc: 0.750863
Early stopping, best iteration is:
[739]   valid_0's auc: 0.751216
```

AUC 점수

```
for base, X in [(base_train, X_train), (base_valid, X_valid), (base_test, X_test)]:
    y_pred = gbm.predict(X, num_iteration=gbm.best_iteration)
    base["score"] = y_pred

print(f'The AUC score on the train set is: {roc_auc_score(base_train["target"], base_train["score"])}')
print(f'The AUC score on the valid set is: {roc_auc_score(base_valid["target"], base_valid["score"])}')
print(f'The AUC score on the test set is: {roc_auc_score(base_test["target"], base_test["score"])}')

The AUC score on the train set is: 0.764122917660593
The AUC score on the valid set is: 0.7512157223309048
The AUC score on the test set is: 0.7483072129459662
```

GINI 계수의 안정성 점수 (통계적 분산 정도)

```
def gini_stability(base, w_fallingrate=88.0, w_resstd=-0.5):
    gini_in_time = base.loc[:, ["WEEK_NUM", "target", "score"]].#
        .sort_values("WEEK_NUM")#
        .groupby("WEEK_NUM")[["target", "score"]]#
        .apply(lambda x: 2*roc_auc_score(x["target"], x["score"])-1).tolist()

    x = np.arange(len(gini_in_time))
    y = gini_in_time
    a, b = np.polyfit(x, y, 1)
    y_hat = a*x + b
    residuals = y - y_hat
    res_std = np.std(residuals)
    avg_gini = np.mean(gini_in_time)
    return avg_gini + w_fallingrate * min(0, a) + w_resstd * res_std

stability_score_train = gini_stability(base_train)
stability_score_valid = gini_stability(base_valid)
stability_score_test = gini_stability(base_test)

print(f'The stability score on the train set is: {stability_score_train}')
print(f'The stability score on the valid set is: {stability_score_valid}')
print(f'The stability score on the test set is: {stability_score_test}')

The stability score on the train set is: 0.4976648127691175
The stability score on the valid set is: 0.4726726686264489
The stability score on the test set is: 0.4583643686935092
```

예측

```
X_submission = data_submission[cols_pred].to_pandas()
X_submission = convert_strings(X_submission)
categorical_cols = X_train.select_dtypes(include=['category']).columns

for col in categorical_cols:
    train_categories = set(X_train[col].cat.categories)
    submission_categories = set(X_submission[col].cat.categories)
    new_categories = submission_categories - train_categories
    X_submission.loc[X_submission[col].isin(new_categories), col] = "Unknown"
    new_dtype = pd.CategoricalDtype(categories=train_categories, ordered=True)
    X_train[col] = X_train[col].astype(new_dtype)
    X_submission[col] = X_submission[col].astype(new_dtype)

y_submission_pred = gbm.predict(X_submission, num_iteration=gbm.best_iteration)

y_submission_pred
array([0.0108652 , 0.0546489 , 0.00741195, 0.00928448, 0.06616796,
       0.01272284, 0.03471827, 0.00326202, 0.06028441, 0.00793056])
```

예측값

```
submission = pd.DataFrame({  
    "case_id": data_submission["case_id"].to_numpy(),  
    "score": y_submission_pred  
}).set_index('case_id')  
submission.to_csv("./submission.csv")
```

| | A | B | C | D | E | F |
|----|---------|----------|---|---|---|---|
| 1 | case_id | score | | | | |
| 2 | 57543 | 0.010865 | | | | |
| 3 | 57549 | 0.054649 | | | | |
| 4 | 57551 | 0.007412 | | | | |
| 5 | 57552 | 0.009284 | | | | |
| 6 | 57569 | 0.066168 | | | | |
| 7 | 57630 | 0.012723 | | | | |
| 8 | 57631 | 0.034718 | | | | |
| 9 | 57632 | 0.003262 | | | | |
| 10 | 57633 | 0.060284 | | | | |
| 11 | 57634 | 0.007931 | | | | |